


One problem, four register wrappers

politeness / length
↓



LLM

model reply

solution code

trailing prose

prose volume grows with politeness

terse appears to beat polite: +15.9 pts, $p = 1.6 \times 10^{-7}$ (artifact)

Naive extractor
keeps the whole reply

✗ SyntaxError
on trailing prose

scored
FAIL

*same completions —
only the extractor differs*

Robust extractor
trims to compilable code

✓ compiles,
run official tests

true
PASS / FAIL

no register effect: 0.0 pts, $p = 1.0$ (ours) / -2.4 pts, $p = 0.27$ (EvalPlus sanitize)